

# An NLP-Based Scotch Whisky Recommender Agent

COMP3004 - Designing Intelligent Agents

Robert Soane

June 2021

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Scotch: Distilled</b>	<b>2</b>
2.1	A brief overview of Scotch . . . . .	2
2.2	Whisky and words . . . . .	2
<b>3</b>	<b>Background and Literature</b>	<b>3</b>
3.1	Language Models . . . . .	3
3.2	Recommender engines . . . . .	4
3.3	AI applications to Whisky . . . . .	4
<b>4</b>	<b>Approach</b>	<b>5</b>
4.1	Requirements . . . . .	5
4.2	The Data . . . . .	5
4.3	Choice of Recommender Method . . . . .	5
4.4	NLP Methods . . . . .	6
<b>5</b>	<b>Implementation</b>	<b>7</b>
5.1	Data Exploration . . . . .	7
5.2	Agent Design . . . . .	11
<b>6</b>	<b>Evaluation</b>	<b>12</b>
6.1	Survey . . . . .	12
6.2	Results . . . . .	12
6.3	Discussion . . . . .	13
<b>7</b>	<b>Conclusion</b>	<b>14</b>
7.1	Suggestions for Future Work . . . . .	14
<b>8</b>	<b>Statement on Ethics</b>	<b>15</b>

## 1 Introduction

Scotch whisky has been produced as early as 1494 by produced by distilling fermented grains to produce a high proof spirit [1, 2]. According to the Scotch Whisky Association, prior to COVID the

Scotch Whisky industry “accounted for 75% of Scotland’s food and drink exports”, and had a year on year growth of 4.4% [3, 4]. With over 130 distilleries, each producing very different flavour profiles, choosing the next whisky to try could be challenging, enthusiast and beginner alike [5, 6].

Limited attempts have been made to apply AI methods to create recommender agents for whiskies, however these focus on customer trend data (which this report argues is not the best strategy), or predominantly use distinct details about whiskies (such as distillery, ABV etc) on which to base their recommendations [7, 8].

This project sought to apply NLP techniques to produce a recommender agent and ascertain whether NLP techniques applied to Whisky tasting notes can power an effective recommender agent.

In section 2 we briefly discuss the production process of Scotch and it’s lexicon. Section 3 discusses the relevant literature in NLP, recommender engines and Scotch. Sections 4 & 5 describe the approach and implementation of the agent, and in section 6 the agent is evaluated experimentally via a questionnaire given to 30 Scotch enthusiasts.

## 2 Scotch: Distilled

This project concerns itself with the specific domain of Scotch. To fully understand the problem at hand, a basic understanding of the drink is required. For this reason, in this section I present a brief overview of Scotch, and a summary of its lexicon. From thereon in whisky-specific terms can be used without explanation.

### 2.1 A brief overview of Scotch

Scotch whisky refers to whisky produced in Scotland fulfilling a set of legal requirements set by the UK Government [9]<sup>1</sup>. Grains are allowed to malt (germinate) to develop their sugars, after which the sugars are extracted to produce a syrup called *wort*. The wort is fermented to produce a sweet hop-free beer. The beer is distilled in *pot stills* to increase the alcohol content significantly to produce *new-make spirit*. This is matured in oak casks for a minimum of 3 years, before the whisky is bottled. At this stage the distiller may choose to dilute the whisky to an ABV of no less than 40% [1, 2].

### 2.2 Whisky and words

When describing whisky, there two important distinctions to be made. Malt versus grain, and single versus blended. *Malt* refers to a whisky wherein the only grain used is *malted barley*, whereas *grain* whisky can be made with mixtures of any grains [10]. *Single* refers to whisky where all whisky in the bottle has come from the same distillery, whereas *blended* whisky is a blend from any number of distilleries. It is important to note that single whiskies are usually themselves blends, but blends of casks from the same distillery [11].

The flavours present in Scotch come from a number of places, and this in turn influences how various whisky’s flavour profiles are described. A descriptor often used to describe whiskies is *peat*. To stop the malting process, the grain is heated, some distilleries (particularly those on the island of Islay) use a peat fire to carry this out. This imparts a smokey flavour onto the grain, which carries through to the end spirit. This smokey flavour is often described as *peated* [1, 12].

The maturation process provides another opportunity to add flavour to the drink. The requirement to age all Scotch in oak casks is resource intensive, and has lead to distilleries purchasing used casks from other drinks manufacturers. Traditionally the sherry industry has supplied used casks to distilleries, and more recently bourbon casks have been used. Any cask which has previously held any drink can be used, be it for the entire maturation process, or at the end such as a *sherry cask finish*. These all add their own flavours to the drink, and this is reflected in whisky tasting notes [1, 13].

---

<sup>1</sup>This report concerns Scotch whisky, and thus *Scotch* and *whisky* are used interchangeably.

## 3 Background and Literature

### 3.1 Language Models

In general, Natural Language Processing (NLP) tasks require a language model of some form or another. Artificial Intelligence (AI) based methods cannot process text in its native unstructured form, but need to convert the raw text to a structured form suitable for the computer to understand. This is often referred to as *embedding*.

The two predominant model types are *syntactic* and *semantic* models. Syntactic methods transform text to a set of ‘symbols’ which carry no inherent meaning, but can be compared across instances in a dataset, whereas semantic methods (such as those described in subsection 3.1.3) retain a general understanding of the text [14].

A dominant syntactic method for transforming unstructured text into a computer-analysable form is the Bag-of-words (BoW) model. The dataset is tokenized (split into individual words), lemmatized (see subsection 3.1.1) and  $k$  keywords are extracted (see subsection 3.1.2) to form our bag of words  $\underline{b} \in \mathbb{R}^k$ . Each document<sup>2</sup> is transformed to a vector  $\underline{v} \in \mathbb{R}^k$  such that  $v_i$  is the frequency of the word  $b_i$  occurring in the document [14, 15, 16].

#### 3.1.1 Stemming and Lemmatization

When dealing with text data, it is not uncommon to have multiple forms of the same word. A syntactic model would view the words ‘cat’ and ‘cats’ as two different discrete symbols. A method is needed to reduce words to a normal form.

Porter proposed an algorithm for removing word suffixes to aim for a normal form, this is called *stemming*. With no semantic understanding, the algorithm searches for specific suffix patterns and removes them until it is unable to [17].

A more semantic approach would be *lemmatization*. Instead of algorithmically removing word endings lemmatization aims to normalise words to a real word root, that is a lemmatizer would reduce the word to the dictionary form of the word [18]. A lemmatizer implementation in Python is the WordNetLemmatizer in the Python Natural Language Tool Kit (NLTK), which queries the WordNet corpus to find the root word [15, 19].

#### 3.1.2 Keyword Extraction

For syntactic methods, keyword extraction (KE) is key. For the purposes of this report, a keyword is a word of particular relevance or importance, and from which we might extract useful information. KE refers to strategies based on which those important words can be ranked, and only the most relevant kept.

**TF-IDF** One such method, is Term Frequency Inverse Document Frequency (TF-IDF). This is commonly used with BoW, and is implemented in Scikit-Learn [20]. TF-IDF is a statistic for scoring a word’s importance based on how frequently it occurs in a document, and how frequently it occurs in the dataset [21].

Scoring as such aims to penalise words that occur too frequently across a document, boosting the scores of words in an individual document for which they have a disproportionately high frequency.

**Graph based KE** Another approach for KE is the use of graph-based ranking methods. These methods model words as nodes on a mathematical network graph<sup>3</sup>. A popular example is the *Rapid*

---

<sup>2</sup>It is common to refer to an instance in a text dataset as a *document*

<sup>3</sup>A graph  $G$  being a set of nodes  $V$  and edges  $E$ . For a brief summary see Rashid Bin Muhammad’s site <http://personal.kent.edu/~rmuhamma/GraphTheory/MyGraphTheory/defEx.htm> [22].

*Automatic Keyword Extraction* (RAKE) algorithm, which splits finds a set of candidate keywords, and models them as a co-occurrence graph.

Each node represents a candidate, each edge co-occurrence, and it's weight the number of co-occurrences. The candidates are then ranked according to frequency and degree (sum of weights) [23].

Beliga et al., survey a wide range of graph based KE techniques, many of which rely on different centrality measures [24]. Onesuch centrality measure which may be useful for this problem is eigencentality [25]. Essentially, eigencentality aims to assign each node as a proportion of the sum of all nodes to which it is connected. Suppose we have a graph, with an adjacency matrix  $A$ , with  $x_i$  being the centrality of the  $i^{th}$  node, we would set  $x_i = \frac{1}{\lambda} A_{ij} x_j$  (using the summation convention). This reduces to the eigenvector equation  $\mathbf{A} \cdot \underline{x} = \lambda \underline{x}$ . This is given with more detail in [26].

### 3.1.3 Word2vec

*Word2vec* is a semantic language model developed by Google. Instead of encoding each word as a discrete symbol as with BoW, word2vec embeddings retain similarity between similar words. This is achieved by training an *Artificial Neural Network* (ANN) to predict the surrounding words for any given word. The weights of the hiddern layer represent probabilities of respective surrounding words. These probability vectors are used as embeddings for each word. As a words embedding now reflects the likely surrounding words, synonyms are mapped to similar vectors. [27, 28, 29]

## 3.2 Recommender engines

Collaborative filtering (CF) is perhaps the more common recommender engine method. CF treats each user as an entity, and provides recommendations to users based on the behaviours of users it deems to be similar [30, 31]. A simple and less abstract example would be an online shopping site recommending product B to someone who has just bought product A on the basis that a significant proportion of shoppers who buy product A go on to purchase B. It produces a simple filter making predictions with no knowledge of either the products, just the user patterns

Content based (CB) recommender engines are the opposite. Instead of focussing on user patterns they make predictions on the basis of specific attributes of each entity being recommended [30, 32]. While such a system may use user details, the main knowledge source is the things being recommended.

## 3.3 AI applications to Whisky

There is a large gap in the research regarding Artificial Intelligence (AI) applications to whisky. Coldevin built a agent based whisky recommender, choosing to use a CB design [8]. He built a system which used specific attributes about a whisky, such as distillery and cask to recommend based on consumer likes or dislikes. Omid-Zohoor and Eghtesadi built another such hybrid (using both CB and CF) agent [7], however again they relied on specific categorical and ratio features for each bottle. An interesting design choice was an agent that predicts on the basis of a users entire profile, and the ratings they give. Perhaps unsurprisingly their CB model performed poorly with users who gave large numbers of reviews. The more reviews the more noise.

I think it is the case that an agent would be far more useful if it took a set of specific user preferences, at a given point in time, and makes a recommendation on that basis. That way, instead of attempting to offer the user a whisky which best matches all of the varied whiskies they like, the user gets a recommendation of the specific style they want at this point in time. This is far closer to coldevin's approach.

In a paper on the subject, Wishart has completed what seems to be the only study of whisky which uses NLP based techniques [33]. Working with industry experts he selected 84 different whiskies and extracted descriptors from their tasting notes. These were coded and used to cluster the whiskies by HCA. These clusters were manually reviewed and evaluated by industry experts. He later proposed his set of 12 flavour dimensions for Scotch whisky [34].

While groundbreaking, Wisharts work differs somewhat from that carried out in this project. Where Wishart aimed to find the key flavours in Scotch, working with industry experts, I aim to produce an agent which exhibits intelligence as an industry expert by suggesting whiskies, based only on their tasting notes.

## 4 Approach

### 4.1 Requirements

The following broad requirements were set out for the agent:

- **Agent and Environment** - The recommender should be an agent acting in an environment. The environment being the contents of the Master of Malt website, and an interface with a user. The agent could be considered as part of a backend of a web app with outputs in a Python dictionary/JSON format.<sup>4</sup>
- **Speed** - The agent should be able to produce a recommendation within a couple of seconds.
- **Customisable** - An end user should be able to filter by price, volume and abv.
- **Updateable** - The agent should be able to automatically update it's database to include new whiskies, and retrain its language model.
- **Input Types** - The agents should recommend based on *likes & dislikes* of whiskies supplied by a user, or from a users written tasting notes.

### 4.2 The Data

In order to build a whisky specific language model (discussed further in subsection 4.4), a large corpus of tasting notes was required. This could also fulfill the requirements of a dataset from which to make predictions.

Product data for a large range of scotch whiskies was scraped from masterofmalt.com<sup>56</sup>, this dataset contains a selection of attributes for each whisky. The name and URL are hashed together using MD5 to provide an ID.

It was observed that whiskies which are discontinued tend to be listed without a price, whereas those which are merely out of stock are listed with a price. For this reason, and for the sake of simplicity, price was taken as an indication of whether a whisky is likely to be in stock. Those without a price were still recorded in the dataset for two reasons; users may wish to make recommendations based on liking or disliking them, and they add to the corpus of tasting notes.

### 4.3 Choice of Recommender Method

While an online spirits shop may use CF recommenders to recommend people viewing a product view another similar product (such as recommending gin to customers who like gin etc.), a CF system may fail for recommending whiskies based on tastes.

It is not unlikely that a whisky drinker may wish to buy two very different whiskies in one order, just to compare them, or they might enjoy a large range of whiskies. Similarly one might only like a large variety of whiskies, and browse and purchase many different styles frequently. It seems somewhat unlikely that shopping habits of whisky geeks would be sufficient to recommend a whisky on the basis of specific tastes.

---

<sup>4</sup>**Note:** The webapp was beyond the scope of this project

<sup>5</sup>masterofmalt.com is a major UK whisky and spirits retailer.

<sup>6</sup>The scraping process is discussed in subsubsection 5.2.2

For that reason, a CB recommender model was chosen to be implemented, based upon tasting note data.

## 4.4 NLP Methods

Word2vec and BoW were both considered as candidate language models. While there are many pretrained models available, these are likely to be unsuitable as briefly explained in subsection 5.1.1.

Word2vec encapsulates a far greater amount of semantic data, however re-training word2vec regularly with new data would be expensive. As a quicker model to train BoW was chosen. TF-IDF, RAKE and an eigencentality ranking measure discussed in subsection 4.4.3 were considered for KE.

### 4.4.1 The Ideal Vector and Similarity

The BoW model maps each input to a vector. To make recommendations, the agent must map the user input to a vector in the same space as the BoW model. This *Ideal Vector* (IV) represents a hypothetical whisky which best represents the user’s input. Cosine similarity can then be used to ascertain which whiskies in the database best match the input. Cosine similarity indicates the angle between vectors [30]. As for  $\underline{u}, \underline{v} \in \mathbb{R}^k$ ,  $\underline{u} \cdot \underline{v} := |\underline{u}||\underline{v}| \cos \theta$  by storing all vectors normalised, this reduces such that the cosine similarity of  $\underline{u}$  &  $\underline{v}$  is simply their scalar product.

$$\begin{pmatrix} d_{11} & d_{12} & \dots & d_{1n} \\ d_{21} & d_{22} & \dots & d_{2n} \\ \dots & \dots & \ddots & \dots \\ d_{m1} & d_{m2} & \dots & d_{mn} \end{pmatrix} \cdot \begin{pmatrix} v_1 \\ v_2 \\ v_3 \\ \dots \\ v_n \end{pmatrix} = \begin{pmatrix} c_1 \\ c_2 \\ c_3 \\ \dots \\ c_n \end{pmatrix} \quad (1)$$

Calculating cosine similarity for a large dataset is relatively straightforward, and this method was implemented in the agent. Consider our dataset of  $m$  whiskies as a matrix  $\mathbf{D} \in \mathbb{R}^{m \times n}$  with each row representing the corresponding whisky’s vector, and our IV  $\underline{v} \in \mathbb{R}^n$ . As demonstrated in equation 1, the product of  $\mathbf{D} \cdot \underline{v} = \underline{c}$  where  $\underline{c}$  is our vector of cosine similarities.

### 4.4.2 Tasting Notes

Quite usefully, as shown in table 1<sup>7</sup>, whisky tasting notes are very keyword dense. This means most words are candidate keywords, however some KE techniques (such as RAKE) are aimed at finding keywords from far less keyword dense text. This must be considered when choosing a KE method.

### 4.4.3 Eigencentality based Rapid Automatic Keyword Extraction (eRAKE)

As discussed in paragraph 3.1.2, co-occurrence graphs can be useful for KE. RAKE is one such method where primitive centrality measures are used to rank nodes. Another such method uses eigencentality. This steps beyond words which themselves have a high co-occurrence and rewards words with significantly weighted edges to words with high co-occurrences. This could be a reasonable compromise for retaining semantic data. If we aren’t keeping full semantic data we are at least stepping beyond merely looking at frequencies, aiming for descriptors with larger amounts of influence across the dataset.

As the graph is undirected, our adjacency matrix is hermitian and thus its eigenvectors relatively simple to find, this can be completed using SciPy’s *eigh()* function [35, 36].

---

<sup>7</sup>All table data from masterofmalt.com

Table 1: A selection of whisky tasting notes from Master of Malt

Whisky	Tasting Notes
Laphroaig 10 Year Old	<p><b>Nose:</b> <i>This opens on big, smoky muscular peat notes. There are spices, and liquorice, as well as a big dose of salt. This whisky has become slightly sweeter in recent years, and it appears beautifully on the nose, amidst the classic iodine/sticking plasters and cool wood smoke we love.</i></p> <p><b>Palate:</b> <i>Seaweed-led, with a hint of vanilla ice cream and more than a whiff of notes from the first aid box (TCP, plasters etc). The oak is big, and muscles its way into the fore as you hold this whisky over your tongue. An upsurge of spices develop – cardamom/black pepper/chilli.</i></p> <p><b>Finish:</b> <i>Big and drying, as the savoury, tarry notes build up with an iodine complexity .</i></p>
Talisker 10 Year Old	<p><b>Nose:</b> <i>A fresh and fragrant nose. Through thick, pungent smoke comes sweet pear and apple peels, with pinches of maritime salt from kippers, seaweed.</i></p> <p><b>Palate:</b> <i>It’s a bonfire of peat crackling with black pepper, with a touch of brine and dry barley. A welcome delivery of orchard fruit provides a delicate and beautiful balance.</i></p> <p><b>Finish:</b> <i>In a long finish, bonfire embers toast malt and crystallise a sugary underlay</i></p>

## 5 Implementation

The implementation was split into the following two broad phases:

- **Data exploration:** Exploring the dataset, and exploring potential prototypical methods in a Jupyter Notebook. Deciding on a chosen method to implement in the agent.
- **Agent design:** Designing an agent in Python based on work from previous phase.

These are discussed in sections 5.1 & 5.2 respectively.

### 5.1 Data Exploration

#### 5.1.1 Effective Lemmatizing

Given the whisky specific lexicon, an interesting problem occurs when trying to use general purpose lemmatizers. In whisky ‘peated’ is a verb describing how the grain was processed, and thus the ‘peaty’ (smokey) flavour has been imparted on the finished spirit. In English, peat refers to a natural fuel made from dead plant matter, and ‘peated’ does not exist [37]. For the purposes of the agent, ‘peated’ and ‘peaty’ should both reduce to ‘peat’, however as they are not considered as verbs in WordNet. For this reason a separate custom whisky lemmatizer was built.

This *WhiskyLemmatizer* was built on top of scikit-learn’s WordNet lemmatizer [20]. By manually creating a dictionary of whisky words and their desired root form, for any input the lemmatizer can first check the dictionary. If the word is not in the dictionary, it can then use the WordNet lemmatizer. The WordNet lemmatizer can be slow as it doesn’t cache results and re-queries them from the WordNet corpus each time. For this reason, every result from WordNet is added to the dictionary.

After experimentation with this *WhiskyLemmatizer*, words in the cache which are mapped to words which further reduce to smaller words were automatically updated to map to the leaf word. A set of stopwords was manually produced in an iterative process based on the lemmatizers outputs.

Table 2: Times of TF-IDF, RAKE and eRAKE with various lemmatizers in seconds.

	TF-IDF	RAKE	eRAKE
Unlemmatized	1247	0.441	-
WordNet Lemmatized	1461	130.2	-
WhiskyLemmatizer	1209	4.947	47.6

**Note:** eRAKE was only applied to the WhiskyLemmatized corpus as the eRAKE implementation included the WhiskyLemmatizer.

### 5.1.2 Comparing KE Strategies

As it is possible that over time the tasting note lexicon may change, the agent should perform KE with each training cycle. KE takes a significant proportion of time for model building. For this reason time and accuracy are of equal importance. It is little use performing KE perfectly if it takes an age.

An adapted implementation of TF-IDF [38] the *rake-nltk* python package [39], and a new implementation eRAKE were applied to the dataset with a range of lemmatizers to extract the top 300 keywords. The methods were timed, and the top 20 keywords recorded. These can be found in tables 2 and 3.

As is clear, the TF-IDF KE took orders of magnitude longer than RAKE and eRAKE. Qualitatively evaluating the keywords extracted by RAKE vs eRAKE, eRAKE produces more useful keywords. This is perhaps unsurprising, as RAKE aims to find keywords from a corpus with both a relatively high frequency of each keyword, and a higher frequency of stopwords. By using it on the corpus of tasting notes it is perhaps being misused. As highlighted in subsection 4.4.2, tasting notes are very feature dense, however different tasting notes have different features characterising them. It is likely that the most frequent features are penalised due to stopword potential. It is interesting to see that WordNet lemmatizing had little impact in terms of which words were extracted.



Table 3: Top 20 keywords from each of TF-IDF, RAKE and eRAKE with various lemmatizers.

Keyword Extraction	Lemmatizer	Keywords
TF-IDF	None	<i>vanilla, quite, juicy, jam, zest, liquorice, crème, waxy, mixed, oak, zesty, smoke, marzipan, drizzle, hazelnut, beeswax, joined, juice, brûlée, box</i>
	Wordnet	<i>vanilla, quite, juicy, jam, zest, liquorice, crème, waxy, mixed, oak, zesty, smoke, marzipan, drizzle, hazelnut, beeswax, joined, juice, brûlée, box</i>
	Whisky	<i>vanilla, zest, jam, quite, juicy, sweet, fruit, waxy, liquorice, crème, smoke, develop, oak, mixed, drizzle, hazelnut, marzipan, join, dry, beeswax</i>
RAKE	None	<i>with, winesky, while, touch, torten, time, theres, saucepan, salty, pan, or, nose, musty, muscular, more, marketplace, little, like, just, its</i>
	Wordnet	<i>with, winesky, while, touch, torten, time, theres, saucepan, salty, pan, or, nose, musty, muscular, more, marketplace, little, like, just, its</i>
	Whisky	-
eRAKE	Whisky	<i>fruit, sweet, spice, oak, vanilla, smoke, honey, malt, chocolate, apple, dry, pepper, orange, cream, butter, fresh, nut, peel, rich, barley</i>

Table 4: Whiskies considered in clustering evaluation.

Highland Park 12 Year Old, Bowmore 15 Year Old, Arran 10 Year Old, Edradour 10 Year Old, Old Pulteney 12 Year Old, Laphroaig 10 Year Old, Ardbeg 10 Year Old, Blair Athol 12 Year Old - Flora and Fauna, Talisker 10 Year Old, GlenAllachie 15 Year Old, Aberlour A'Bunadh Batch 68

### 5.1.3 Clustering

Table 5: Clustering of whiskies from each BoW model.

KE	Cluster tures	Fea- tures	Whiskies	KE	Cluster tures	Fea- tures	Whiskies
TF-IDF	<i>spice, vanilla, sweet, malt, honey, fruit, malt, spice, oak, vanilla, smoke, malt, chocolate, malt, sherry</i>		Highland Park, Bowmore Arran Edradour Old Pulteney Laphroaig, Ardbeg, Blair Athol, Talisker GlenAllachie, Aberlour A'Bunadh	TF-IDF*	<i>fruit, malt, spice, malt, fruit, oak, fruit, malt, spice, oak, vanilla, smoke, peat, malt, sherry, malt, fruit</i>		Highland Park, Bowmore Arran - Old Pulteney Laphroaig, Ardbeg, Talisker GlenAllachie, Blair Athol, Aberlour A'Bunadh, Edradour
RAKE	<i>musty, muscular, fire, little, musty, fire, little, like, musty, time, like, fire, like, muscular, musty, little, time, like</i>		Laphroaig, Ardbeg, Highland Park, Old Pulteney, GlenAllachie, Blair Athol, Bowmore, Aberlour A'Bunadh, Edradour - Talisker Arran -	eRAKE	<i>malt, vanilla, sweet, sherry, malt, fruit, malt, fruit, oak, fruit, malt, spice, smoke, peat, vanilla, oak, malt, vanilla</i>		Highland Park, Bowmore GlenAllachie, Blair Athol, Aberlour A'Bunadh, Edradour Arran - Laphroaig, Ardbeg, Talisker Old Pulteney

TF-IDF and RAKE refer to these extractions applied both WordNet and unlemmatized. They produced the same results. TF-IDF\* is TF-IDF used with WhiskyLemmatizer

Cluster Features refers to the three most prominent features at the centers of the cluster.

For the purposes of sanity checking, and ensuring sufficient information is retained in each BoW model, k-means clustering was applied on a BoW model based on each set of keywords. The clusters of the whiskies in Table 4 were considered<sup>8</sup> and the corresponding clusters are shown in Table 5.

<sup>8</sup>I chose these whiskies as I have enough basic knowledge of them to qualitatively approximately evaluate the sensibleness of the clustering and help make a decision. This is hardly a rigorous approach, and future work would need a more rigorous evaluation at this stage. This faster approach was used due to time constraints.

There is little difference between TF-IDF and TF-IDF\*, apart from Edradour being clustered with GlenAllachie and Aberlour (two heavily sherried expressions) in TF-IDF\*, and Blair Athol. Blair Athol is a relatively sherried expression, and thus appearing in a very smoke and peat heavy cluster (Laphroaig, Ardbeg, Talisker) seems strange. Its placement in TF-IDF\* seems far more sensible.

When considering the tasting notes, Blair Athol is described as “*Nutty with sherried notes. Gentle peat. Crisp. ... Peat smoke, syrup ...*” [40]. This highlights a limitation of the BoW model - and it’s applications to this problem. by mentioning peat and smoke three times, Blair Athol was grouped with similar peat and smoke heavy whiskys.

The RAKE clusters are clearly nonsense, however this isn’t surprising when considering the main features. The eRAKE clusters are very similar to those in TF-IDF\*. On the basis of this, and the data in subsection 5.1.2, eRAKE was chosen to move forward with.

## 5.2 Agent Design

The agent was designed as a single Python class, with a few helper classes and functions. All key features of the agent can be accessed through a small subset of methods. While the agent isn’t directly useable by a user, all inputs and outputs use Python dictionaries. While this may seem strange as the methods give less information regarding their expected input/output, this is with the view that a web frontend could be designed to send data in JSON formats.

### 5.2.1 The Database

An SQLite database is used to store all whisky data and models. This is implemented directly in Python and allows multiple agents to access the data at the same time. This also allows one agent to update the models and all other agents can use the up to date model. It was observed that SQLite appears to run faster than Pandas, however Pandas<sup>9</sup> is still used for some data manipulations once data is loaded from the database.

When loaded the agent checks if there’s a *.DB* directory, and if there isn’t it creates it. Using the pre-existing *scotch.csv* file a database is created with a table for each product, model and review.

### 5.2.2 Web Scraping

The initial dataset was collected using a roughly hacked together script using Python and BeautifulSoup [42]. The agent was designed using this dataset of approximately 14,000 whiskies. As per the requirements, and to ensure the agent’s autonomy, a method was written to fetch new whiskies. While it is not recommended that it is used, if the initial dataset is not present in the agents root folder, it will automatically fetch all data from Master of Malt before creating the database.

Master of Malt list new whiskies on a page of their website. This page is parsed and each listing is checked to confirm it is Scotch. When each ID is created, it is checked against those already in the database. If three consecutive listings are already in the database, the agent stops assuming all new products have been included. The reason for this is sometimes one or two re-stocked whiskies are listed in succession. Setting three as the threshold reduces the risk of stopping prematurely.

While this function could be setup to run periodically, this hasn’t been implemented to avoid unnecessarily hammering the Master of Malt server.

### 5.2.3 Model Training

The eRAKE method is used to lemmatize all tasting notes, and extract 300 keywords for each model. These are then vectorised and the dataset is transformed to a matrix  $\mathbf{M} \in \mathbb{R}^{m \times n}$ , with  $m$  whiskies,  $n$  keywords, and each row normalised. Each  $\mathbf{M}$  is stored with each whiskies corresponding ID in a table. Four models are created, described in Table 6.

<sup>9</sup>A python package for manipulating data tables [41]

Table 6: Description of models created for agent.

Model	Description
Nose	Model based purely on keywords extracted from <i>nose</i> tasting notes
Palate	Model based purely on keywords extracted from <i>palate</i> tasting notes
Finish	Model based purely on keywords extracted from <i>finish</i> tasting notes
General	Model based on keywords extracted from all tasting notes. Vectorising description as well as tasting notes for each whisky. This reflects some whiskies being listed without tasting notes, but taste indications in main description.

#### 5.2.4 Producing Recommendations

As discussed in section 4.4.1, cosine similarity and matrix algebra is used to produce recommendations. Where recommendations are made based on more than one model, the mean of the cosine similarities is used. A problem encountered with Pandas was that it can be quite slow due to it’s single-threaded nature, especially when converting a long list of data entries into a dataframe. To minimise this effect, the users filtering input (such as by price, ABV etc.) are used to select a get the set of all IDs for from which a recommendation can be made. Only the models for these whiskies are queried from the database.

#### 5.2.5 Dream Dram

An interesting feature implemented is the *Dream Dram* (DD) recommender. This takes unstructured text describing a ‘dream’ whisky and makes a recommendation on that basis. This works in much the same way as the other recommendations, however the IV is generated on the basis of the text input instead of by querying for specific whiskies. This option could potentially be incorporated into a Whisky chat bot at a later date.

## 6 Evaluation

### 6.1 Survey

To evaluate the performance of the agent, a sample dataset of potential user inputs and agent outputs was produced. A description of each input can be found in Table 7. Random sampling from the dataset was used to produce a baseline. These were used in a survey which was given to 30 whisky enthusiasts to complete<sup>10</sup>. They were asked to rate each recommendation on the basis of it’s corresponding user input. While the agent recommends 10 whiskies by default, only 3 recommendations were give from the baseline and agent to avoid making the survey too cumbersome.

The participants were given all information available to the agent for each output including description and tasting notes (and some information the agent does not use such as ABV, Price etc), however they were instructed not to factor price in their ratings as price parameters were set by the hypothetical user. There was also an option to leave text feedback.

### 6.2 Results

As can be seen in figures 1 and 2, most recommendations were rated higher than the mean baseline.

A paired one-tailed t-test was conducted on the baseline scores and the corresponding recommender scores for each sample input. There was a significant increase in scores from the model ( $M = 6.28$ ,

<sup>10</sup>See section 8 regarding ethics.

Table 7: Description of inputs in evaluation dataset.

Input Reference	Rationale
ATN1	Replicating a user who has tried and developed tastes for a variety of whiskies available at supermarkets, but hasn't tried much beyond.
ATN2	Replicating a significant partiality towards heavily peated whiskies.
ATN3	Replicating an enjoyment of both peated and sherried whiskies.
ATN4	A user with niche and specific whisky tastes.
GC	Producing recommendations based on general inputs without considering specific tasting notes.
DD1	Dream Dram recommendation from a very peat heavy input.
DD2	Dream Dram recommendation describing a very oily and fruity whisky.

$SD = 1.13$ ) compared with the baseline ( $M = 4.71$ ,  $SD = 1.15$ ),  $t(12) = 2.22$ ,  $p < 0.05$ . There was a mean increase of 1.57 points.




Figure 1: Boxplots of participant ratings for each recommendation, grouped by sample input. The grey line indicates the mean baseline rating.

### 6.3 Discussion

Despite the promising results, it is clear that some recommendations received a large range of ratings. The small samples used, both of enthusiasts and inputs/recommendations, are far from ideal.

Further on reading the open ended free-text responses, it's clear that some participants weren't confident in their knowledge of whisky. Similarly some participants hadn't understood the instructions, for example one participant stated that they had penalised recommendations with high prices despite the instructions requesting prices aren't considered.

A few interesting points were made which are worth considering. One participant suggested that the sample input for ATN1 was contradictory, perhaps explaining its relatively low score. This suggests that (as intended) this recommender method is not ideal for taking a users entire taste profile and




Figure 2: Boxplots of participant ratings for each sample input. The grey line indicates the mean baseline rating.

providing recommendations, rather recommending a whisky to try based on similarity between a small number of tried whiskies. It would be interesting to investigate whether such a system is possible in the specific case of whisky.

Many users pointed out that many of the recommendations were very niche, suggesting an extra filter could be incorporated allowing a user to filter out single-cask/independent bottles. This also could have had an impact on accuracy of participants scores.

## 7 Conclusion

This project sought to investigate whether NLP techniques could be applied to free text tasting notes to produce an effective Scotch whisky recommender engine.

Acting on an environment consisting of a retailers website and accepting JSON input, an autonomous agent was produced with capacity to maintain an updated language model for Scotch whisky. This model can reflect changes in the Scotch lexicon, and be applied effectively to recommend whisky on the basis of taste. This could have significant positive consequences for whisky buyers.

Due to the sparsity of the field, there is a wealth of space for exploration, with a few suggestions listed in subsection 7.1. Despite the promising early results, a more comprehensive study is needed, ideally with a larger sample size and/or industry experts.

### 7.1 Suggestions for Future Work

The following suggestions are made for future work and research to build upon this project.

- Developement of a front end UI for the agent.
- Further comparisons and reasearch into good KE methods for whisky tasting notes.
- Investigations into semantic language models and alternative similarity measures.

- Work on whisky clustering to research flavour profile metrics similar to those used in the wine industry.
- Further evaluation of effectiveness of recommender agent.
- Work with industry experts to evaluate and improve the recommender agent.

## 8 Statement on Ethics

*I have followed the guidance given in the module about research ethics in the user study in my coursework for this module.*

## References

- [1] K. Jacques, T. Lyons, and D. Kelsall, *The Alcohol Textbook 4th edition*, 4th ed. Nottingham University Press, 2003.
- [2] M. Pyke, “THE MANUFACTURE OF SCOTCH GRAIN WHISKY,” *The Distillers Company Ltd., Glenochil Research Station, Menstrie, Clackmannanshire, Scotland*, vol. 71, pp. 209–218, 1965.
- [3] “Facts & figures.” [Online]. Available: <https://www.scotch-whisky.org.uk/insights/facts-figures/>
- [4] “Scotch whisky export figures 2019.” [Online]. Available: <https://www.scotch-whisky.org.uk/newsroom/scotch-whisky-exports-surge-amidst-backdrop-of-tariff-uncertainty/>
- [5] “How many whisky distilleries are in scotland?” Oct 2020. [Online]. Available: <https://whiskytastingcompany.com/blogs/news/how-many-whisky-distilleries-are-in-scotland>
- [6] T. Powell, “The beginner’s guide to scotch whisky,” Jan 2021. [Online]. Available: <https://foodism.co.uk/guides/scotch-whisky-regions-guide/>
- [7] A. Omid-zohoor and A. Eghtesadi, “Whisky Recommender.”
- [8] T. M. Coldevin, “Building an Multi-Agent Whisky Recommender System,” no. February, 2005.
- [9] “The scotch whisky regulations 2009,” *Legislation.gov.uk*, 2009. [Online]. Available: <https://www.legislation.gov.uk/ukxi/2009/2890/contents/made>
- [10] P. Valaer, “Scotch Whisky,” *Industrial and Engineering Chemistry*, vol. 32, no. 7, pp. 935–943, 1940.
- [11] B. C. Smith, C. Sester, J. Ballester, and O. Deroy, “The perceptual categorisation of blended and single malt Scotch whiskies,” *Flavour*, vol. 6, no. 1, pp. 1–9, 2017.
- [12] G. N. Bathgate, “The influence of malt and wort processing on spirit character: the lost styles of Scotch malt whisky,” *Journal of the Institute of Brewing*, vol. 125, no. 2, pp. 200–213, 2019.
- [13] J. Mosedale and J.-L. Puech, “Wood maturation of distilled beverages,” *Trends in Food Science & Technology*, vol. 9, no. 3, pp. 95–101, mar 1998. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S0924224498000247>
- [14] E. Cambria and B. White, “Jumping NLP Curves: A Review of Natural Language Processing Research [Review Article],” *IEEE Computational Intelligence Magazine*, vol. 9, no. 2, pp. 48–57, may 2014. [Online]. Available: <http://ieeexplore.ieee.org/document/6786458/>

- [15] E. L. Steven Bird, Ewan Klein, *Natural Language Processing with Python*. O'Reilly Media Inc, 2009.
- [16] Y. Zhang, R. Jin, and Z. H. Zhou, "Understanding bag-of-words model: A statistical framework," *International Journal of Machine Learning and Cybernetics*, vol. 1, no. 1-4, pp. 43–52, 2010.
- [17] M. Porter, "An algorithm for suffix stripping," *Program*, vol. 14, no. 3, pp. 130–137, mar 1980. [Online]. Available: <https://www.emerald.com/insight/content/doi/10.1108/eb046814/full/html>
- [18] K. Jayakodi, M. Bandara, I. Perera, and D. Meedeniya, "WordNet and cosine similarity based classifier of exam questions using bloom's taxonomy," *International Journal of Emerging Technologies in Learning*, vol. 11, no. 4, pp. 142–149, 2016.
- [19] "What is wordnet?" 2010. [Online]. Available: <https://wordnet.princeton.edu/>
- [20] P. Fabian, G. Varoquaux, A. Gramfort, M. Vincent, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research*, vol. 12, no. 85, pp. 2825–2830, 2011.
- [21] J. Ramos, "Using TF-IDF to Determine Word Relevance in Document Queries," *Proceedings of the first instructional conference on machine learning*, vol. 242, no. 1, pp. 29–48, 2003.
- [22] R. B. Muhammad, "Graph theory: Definitions and examples." [Online]. Available: <http://personal.kent.edu/~rmuhamma/GraphTheory/MyGraphTheory/defEx.htm>
- [23] S. Rose, D. Engel, N. Cramer, and W. Cowley, "Automatic keyword extraction," *Text Mining: Applications and Theory*, pp. 1—277, 2010.
- [24] S. Beliga, A. Mestrovic, and S. Martincic-Ipsic, "An Overview of Graph-Based Keyword Extraction Methods and Approaches," *Journal of Information and Organizational Sciences*, vol. 39, no. 1, 2015. [Online]. Available: <https://hrcak.srce.hr/140857>
- [25] P. Bonacich, "Some unique properties of eigenvector centrality," *Social Networks*, vol. 29, no. 4, pp. 555–564, oct 2007. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S0378873307000342>
- [26] M. E. J. Newman, "Mathematics of networks," *Networks*, pp. 109–167, 2010.
- [27] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," in *1st International Conference on Learning Representations, ICLR 2013 - Workshop Track Proceedings*, 2013, pp. 1–12.
- [28] C. McCormick, "Word2Vec Tutorial - The Skip-Gram Model," pp. 1–39, 2017.
- [29] Q. Liu, M. J. Kusner, and P. Blunsom, "A Survey on Contextual Embeddings," 2020. [Online]. Available: <http://arxiv.org/abs/2003.07278>
- [30] P. Melville and V. Sindhvani, "Encyclopaedia of Machine Learning: Recommender Systems," *Encyclopaedia of Machine Learning*, pp. 829–838, 2010. [Online]. Available: [https://link.springer.com/content/pdf/10.1007/978-1-4899-7687-1\\_964.pdf%0Ahttps://link.springer.com/referenceworkentry/10.1007%2F978-0-387-30164-8\\_705%0Ahttps://link.springer.com/book/10.1007/978-0-387-30164-8%0A%0Ahttp://vikas.sindhvani.org/recommender.p](https://link.springer.com/content/pdf/10.1007/978-1-4899-7687-1_964.pdf%0Ahttps://link.springer.com/referenceworkentry/10.1007%2F978-0-387-30164-8_705%0Ahttps://link.springer.com/book/10.1007/978-0-387-30164-8%0A%0Ahttp://vikas.sindhvani.org/recommender.p)
- [31] J. Herlocker, J. Konstan, and J. Reidl, "Explaining Collaborative Filtering Recommendations," Mineapolis, 2000. [Online]. Available: <http://www.grouplens.org>



- [32] R. J. Mooney and L. Roy, “Content-based book recommending using learning for text categorization,” *Proceedings of the ACM International Conference on Digital Libraries*, pp. 195–204, 2000.
- [33] D. Wishart, “Classification of Single Malt Whiskies,” 2000, pp. 89–94. [Online]. Available: [http://link.springer.com/10.1007/978-3-642-59789-3\\_{\\_}14](http://link.springer.com/10.1007/978-3-642-59789-3_{_}14)
- [34] —, “The flavour of whisky,” *Significance*, vol. 6, no. 1, pp. 20–26, mar 2009. [Online]. Available: <http://doi.wiley.com/10.1111/j.1740-9713.2009.00337.x>
- [35] M. Hubbard, “Lecture: Best approximation (2)/eigenvalues and eigenvectors (1), 17/03/2020,” *MATH3036-1-UNUK-SPR-1920 Scientific Computation and Numerical Analysis*, Mar 2020.
- [36] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. Fernández del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, and T. E. Oliphant, “Array programming with NumPy,” *Nature*, vol. 585, p. 357–362, 2020.
- [37] L. Jenner, “peat,” 2019.
- [38] P. Prakash, “Keyword extraction: Keyword extraction in python,” Dec 2020. [Online]. Available: <https://www.analyticsvidhya.com/blog/2020/11/words-that-matter-a-simple-guide-to-keyword-extraction-in-python/>
- [39] V. B. Sharmer, “rake-nltk,” Jun 2018. [Online]. Available: <https://pypi.org/project/rake-nltk/>
- [40] “Blair athol 12 year old - flora and fauna.” [Online]. Available: <https://www.masterofmalt.com/whiskies/blair-athol-12-year-old-whisky/>
- [41] T. pandas development team, “pandas-dev/pandas: Pandas,” Feb. 2020. [Online]. Available: <https://doi.org/10.5281/zenodo.3509134>
- [42] L. Richardson, “Beautiful soup documentation,” *April*, 2007.